

Teaching HPC with a small Raspberry Pi Cluster

Colin Sauze <cos@aber.ac.uk>

Research Software Engineer
Super Computing Wales Project
Aberystwyth University



SUPERCOMPUTING WALES
UWCHGYFRIFIADURA CYMRU

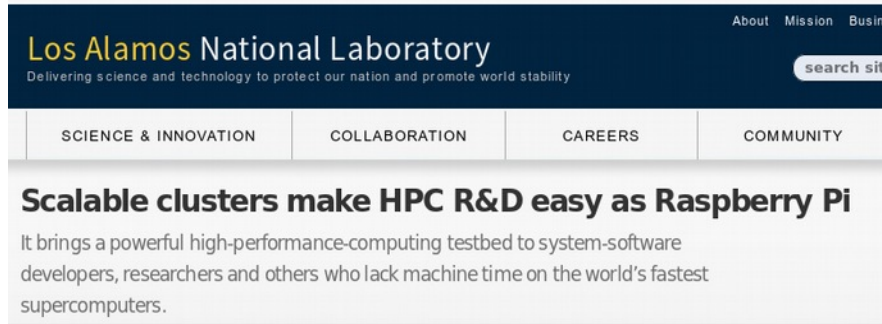


PRIFYSGOL
ABERYSTWYTH
UNIVERSITY


Overview

- Inspirations
- Why teach HPC with Raspberry Pi?
- My Raspberry Pi cluster
- Experiences from teaching
- Future Work

Inspiration #1: Los Alamos National Laboratory



November 13, 2017




- 750 node cluster
- Test system for software development
- Avoid tying up the real cluster

Inspiration #2: Wee Archie/Archlet



- EPCC's Raspberry Pi Cluster
- Archie: 18x Raspberry Pi 2's (4 cores each)
- Archlet: smaller 4 or 5 node clusters.
- Used for outreach demos.
- Setup instructions:
https://github.com/EPCCed/wee_archlet

Inspiration #3: Swansea's Raspberry Pi Cluster



- 16x Raspberry Pi 3s
- CFD demo using a Kinect sensor
- Demoed at the Swansea Festival of Science last year
- More about this in Ed Bennett's talk on 12pm tomorrow

Why Teach with a Raspberry Pi cluster?

- Avoid loading real clusters doing actual research
 - Less fear from learners that they might break something
- Resource limits more apparent
- More control over the environment
- Hardware less abstract
- No need to have accounts on a real HPC

My Cluster



- “Tweety Pi”
 - 10x Raspberry Pi model B version 1s
 - 1x Raspberry Pi 3 as head/login node
 - Raspbian Stretch
 - Slurm job manager
- Head node acts as WiFi access point
 - Internet via phone or laptop
- Used a British Science Week 2019
 - Simple Pi with Monte Carlo methods demo

Teaching Materials

- Based on Introduction to HPC with Super Computing Wales carpentry style lesson:
 - What is an HPC?
 - Logging in
 - Filesystems and transferring data
 - Submitting/monitoring jobs with Slurm
 - Profiling
 - Parallelising code, Amdahl's law
 - MPI
 - HPC Best Practice

Experiences from Teaching

- Ran two courses, part of STFC summer school
- 15 people each time
 - 1st time using HPC for many
 - Most had some Unix experience
- Subset of Super Computing Wales introduction to HPC carpentry lesson
- Slurm issues during first workshop
- Worries about WiFi

Technical Problems

- Compute nodes run root filesystem over NFS from the head node
- Created accounts on the head node
- Copied /etc/passwd to the NFS share
- Original files were different, extra Slurm related system user on the nodes
- Broke Slurm :(
- Fixed for second group

Feedback

- Very Positive
- A lot seemed to enjoy playing around with SSH/SCP
 - First time using a remote shell for some
 - Others more adventurous than they might have been on a real HPC
- Main complaint was lack of time (only 1.5 hours)
 - Only got as far as covering basic job submission
 - Quick theoretical run through of MPI and Amdahl's law
 - Probably have 3-4 hours of material
- Queuing became very apparent
 - 10 nodes, 15 users
 - “watch queue” running on screen during practical parts

Future Work

- More complete setup instructions (link on next slide)
- Script based install
- CentOS/Open HPC stack instead of Raspbian
- Use for full intro to SCW course
 - Hybrid of Raspberry Pi and real cluster
- Public engagement demo which focuses on our research
 - Analysing satellite imagery
 - Simulate the monsters from MonsterLab (<https://monster-lab.org/>)

More Information

- My setup instructions - https://github.com/colinsauze/pi_cluster
- Teaching material - <https://github.com/SCW-Aberystwyth/Introduction-to-HPC-with-RaspberryPi>
- Email me: cos@aber.ac.uk